

Instruction Set Architecture (ISA)

In Systems Organisation 1, we built a computing system from the bottom up, focusing on the hardware, and manipulation of bits using logic.

We will now look at the same problem from a software perspective.

Outline

Machine instructions and program execution

An instruction in the machine corresponds to a pathway through the machine.

E.g. a pathway between the operands of an ADD operation and the piece of hardware that does the addition, and then from that hardware to a place to store the result.

Addressing methods for data operands

We may want to refer to data in memory or a register, or data that doesn't yet exist when the program starts.

The addressing modes allow us to access data that's in different places in our machine.

Assembly-language representation for instructions, data and programs

User-friendly labels and languages.

Stacks and subroutines

A subroutine is a small sub-program, e.g. function calls. The stack allows us to do this.

Things can be computed without a stack, but it's needed practically for programming.

Memory organisation

Memory is made up of millions of cells. Each one holds a bit, 0 or 1.

A word is a group of n bits. Word length can be 16 to 64 bits.

Memory, then, is a collection of consecutive words of the size specified by the word length.

Word and byte encoding

A common word length is 32 bits. This word can store a 32-bit signed integer or four 8-bit bytes, e.g. ASCII characters.

For 32-bit integer encoding, bit b_{31} is the sign bit.

Words in memory may store data or machine instructions for a program.

Each machine instruction may require one or more consecutive words for encoding.

Addresses for memory locations

To store or retrieve items of information, each memory location has a distinct address.

Numbers 0 to $(2^k - 1)$ are used as addresses for successive locations in the memory.

The 2^k locations make up the address space.

The memory size is set by k , the number of address bits. When $k = 32$, we have about 4 billion locations.

Byte addressability

Bytes are always 8 bits, but word length may range from 16 to 64 bits.

It's impractical to assign an address to each bit, so instead the memory is byte-addressable, with each byte having an address.

Byte locations have addresses 0, 1, 2, ... Assuming a word length of 32 bits, word locations have addresses 0, 4, 8, ...

Big- and little-endian addressing

Two ways to assign byte addresses across words.

Big-endian assigns lower addresses to more significant (leftmost) bytes of the word. Little-endian goes the opposite way.

Commercial computers use either approach, and some can support both.

The addresses for 32-bit words are still 0, 4, 8, ...

Word alignment

Numbers of bytes per word is normally a power of 2.

Word locations have aligned addresses if they begin at byte addresses that are multiples of the number of bytes in a word.

Some computers permit unaligned addresses.