

Introduction to Java (cs2514)

Lecture 17: Bloom Filters

M. R. C. van Dongen

March 17, 2017

- This lecture studies *bitmaps* and *Bloom Filters*.
- Bitmaps are space-optimal data structures for representing sets.
- Candidate set members come from an index set $\{0, \dots, n-1\}$.
- They can answer simple questions about the set:
 - Is a candidate member in the set?
- *Bloom Filters* are also used to represent sets.
- Here the members come from a key set $\{k_0, \dots, k_{n-1}\}$.
- The answers from Bloom Filters are *probabilistic*:
 - Not all answers are correct with 100% probability.
- When n gets large, we cannot represent such sets as bitmaps.
- However, Bloom Filters for such sets are small.
- They require only a few bits per element.

Bitmaps

- A *bitmap* represents the members of a given index set.
- Here an index set is a set of the form $\{0, \dots, n-1\}$.
- A bitmap may be represented using a boolean array.
- Java does not prescribe how represent a boolean.
- If a boolean is represented as an `int` then this wastes memory.
- Especially if the index size is large.
- In practice bitmaps are often represented as `int` arrays.

Basic Representation

`Bitmap(int capacity)`: Create a bitmap.

`add(int index)`: Add a given index.

`remove(int index)`: Remove a given index.

`size()`: What is the cardinality of the bitmap?

`contains(int index)`: Does the bitmap contains an index?

Implementation

Java

```
public class Bitmap {
    private final boolean[] bits;
    private int size;

    public Bitmap( int capacity ) {
        bits = new boolean[ capacity ];
        size = 0;
    }

    public void add( int index ) {
        size += (bits[ index ]) ? 0 : 1;
        bits[ index ] = true;
    }

    public void remove( int index ) {
        size -= (bits[ index ]) ? 1 : 0;
        bits[ index ] = false;
    }

    public int size( ) {
        return size;
    }

    public boolean contains( int index ) {
        return bits[ index ];
    }
}
```



Implementation

Java

```
public class Bitmap {
    private final boolean[] bits;
    private int size;

    public Bitmap( int capacity ) {
        bits = new boolean[ capacity ];
        size = 0;
    }

    public void add( int index ) {
        size += (bits[ index ]) ? 0 : 1;
        bits[ index ] = true;
    }

    public void remove( int index ) {
        size -= (bits[ index ]) ? 1 : 0;
        bits[ index ] = false;
    }

    public int size( ) {
        return size;
    }

    public boolean contains( int index ) {
        return bits[ index ];
    }
}
```



Implementation

Java

```
public class Bitmap {
    private final boolean[] bits;
    private int size;

    public Bitmap( int capacity ) {
        bits = new boolean[ capacity ];
        size = 0;
    }

    public void add( int index ) {
        size += (bits[ index ]) ? 0 : 1;
        bits[ index ] = true;
    }

    public void remove( int index ) {
        size -= (bits[ index ]) ? 1 : 0;
        bits[ index ] = false;
    }

    public int size( ) {
        return size;
    }

    public boolean contains( int index ) {
        return bits[ index ];
    }
}
```



Implementation

Java

```
public class Bitmap {
    private final boolean[] bits;
    private int size;

    public Bitmap( int capacity ) {
        bits = new boolean[ capacity ];
        size = 0;
    }

    public void add( int index ) {
        size += (bits[ index ]) ? 0 : 1;
        bits[ index ] = true;
    }

    public void remove( int index ) {
        size -= (bits[ index ]) ? 1 : 0;
        bits[ index ] = false;
    }

    public int size( ) {
        return size;
    }

    public boolean contains( int index ) {
        return bits[ index ];
    }
}
```



Implementation

Java

```
public class Bitmap {
    private final boolean[] bits;
    private int size;

    public Bitmap( int capacity ) {
        bits = new boolean[ capacity ];
        size = 0;
    }

    public void add( int index ) {
        size += (bits[ index ]) ? 0 : 1;
        bits[ index ] = true;
    }

    public void remove( int index ) {
        size -= (bits[ index ]) ? 1 : 0;
        bits[ index ] = false;
    }

    public int size( ) {
        return size;
    }

    public boolean contains( int index ) {
        return bits[ index ];
    }
}
```



Simulation

Java


```
Bitmap set = new Bitmap( 10 );  
set.add( 4 );  
set.add( 8 );  
set.add( 1 );  
set.remove( 8 );  
set.add( 4 );
```

Simulation

Java

```
Bitmap set = new Bitmap( 10 );  
set.add( 4 );  
set.add( 8 );  
set.add( 1 );  
set.remove( 8 );  
set.add( 4 );
```

size:0



Simulation

Java

```
Bitmap set = new Bitmap( 10 );  
set.add( 4 );  
set.add( 8 );  
set.add( 1 );  
set.remove( 8 );  
set.add( 4 );
```

size:1

bits:	0	1	2	3	1	5	6	7	8	9
	0	1	2	3	4	5	6	7	8	9

Simulation

Java

```
Bitmap set = new Bitmap( 10 );  
set.add( 4 );  
set.add( 8 );  
set.add( 1 );  
set.remove( 8 );  
set.add( 4 );
```

size:2

bits:

0	0	0	0	1	0	0	0	1	0
0	1	2	3	4	5	6	7	8	9

Simulation

Java

```
Bitmap set = new Bitmap( 10 );  
set.add( 4 );  
set.add( 8 );  
set.add( 1 );  
set.remove( 8 );  
set.add( 4 );
```

size:3

bits:

0	1	0	0	1	0	0	0	1	0
0	1	2	3	4	5	6	7	8	9

Simulation

Java

```
Bitmap set = new Bitmap( 10 );  
set.add( 4 );  
set.add( 8 );  
set.add( 1 );  
set.remove( 8 );  
set.add( 4 );
```

size:2

bits:

0	1	0	0	1	0	0	0	0	0
0	1	2	3	4	5	6	7	8	9

Simulation

Java

```
Bitmap set = new Bitmap( 10 );  
set.add( 4 );  
set.add( 8 );  
set.add( 1 );  
set.remove( 8 );  
set.add( 4 );
```

size:2

bits:

0	1	0	0	1	0	0	0	0	0
0	1	2	3	4	5	6	7	8	9

Using int Representation

- Remember a Java int is implemented with 32 bits.
- We may represent a bitmap with capacity of n using $\lceil n/32 \rceil$ ints.
- The Integer class defines the number of bits in an int:
 - Integer.SIZE.

Java

```
private final int[] bits;  
private int size;  
public Bitmap( int capacity ) {  
    bits = new int[ (capacity + Integer.SIZE - 1) / Integer.SIZE ];  
    size = 0;  
}
```

- We use bit n to represent index n .

Using int Representation

- Remember a Java int is implemented with 32 bits.
- We may represent a bitmap with capacity of n using $\lceil n/32 \rceil$ ints.
- The Integer class defines the number of bits in an int:
 - Integer.SIZE.

Java

```
private final int[] bits;  
private int size;  
public Bitmap( int capacity ) {  
    bits = new int[ (capacity + Integer.SIZE - 1) / Integer.SIZE ];  
    size = 0;  
}
```

- We use bit n to represent index n .
- But how do we turn the bits on and off?

Manipulating Bits

- Implementing bit operations may be done with arithmetic.
- For example, let `bits` be an `int`.
- Then $(bits / 1) \% 2$ is the zeroth (least significant) bit.
- Then $(bits / 2) \% 2$ is the first bit.
- Then $(bits / 4) \% 2$ is the second bit.
- Then $(bits / 2^n) \% 2$ is the n th bit.

Manipulating Bits

- Implementing bit operations may be done with arithmetic.
- For example, let `bits` be an `int`.
- Then $(bits / 1) \% 2$ is the zeroth (least significant) bit.
- Then $(bits / 2) \% 2$ is the first bit.
- Then $(bits / 4) \% 2$ is the second bit.
- Then $(bits / 2^n) \% 2$ is the n th bit.
- $bits - (((bits / 2^n) \% 2) * 2^n)$ turns the n th bit off.

Manipulating Bits

- Implementing bit operations may be done with arithmetic.
- For example, let `bits` be an `int`.
- Then $(bits / 1) \% 2$ is the zeroth (least significant) bit.
- Then $(bits / 2) \% 2$ is the first bit.
- Then $(bits / 4) \% 2$ is the second bit.
- Then $(bits / 2^n) \% 2$ is the n th bit.
- $bits - (((bits / 2^n) \% 2) * 2^n)$ turns the n th bit off.
- $bits + ((1 - (bits / 2^n) \% 2) * 2^n)$ turns the bit on.

Manipulating Bits (Continued)

Java

```
final static int SHIFT = Integer.numberOfTrailingZeros( Integer.SIZE );  
final static int MASK  = Integer.SIZE - 1;  
public boolean contains( int index ) {  
    return (bits[ index >> SHIFT ] & (1 << (index & MASK))) != 0;  
}
```

Bitwise Operations: Not Examinable

Operator	Example	Result	Description
<code>a >> b</code>	<code>32 >> 4</code>	2	Shift a right by b bits. (Keep sign.)
<code>a >>> b</code>	<code>-1 >>> 30</code>	3	Shift a right by b bits. (Reset sign.)
<code>a << b</code>	<code>3 << 3</code>	24	Shift a left by b bits.
<code>a & b</code>	<code>5 & 3</code>	1	Bitwise and of a and b.
<code>a b</code>	<code>5 3</code>	7	Bitwise or of a and b.
<code>a ^ b</code>	<code>5 ^ 3</code>	6	Bitwise exclusive or of a and b.
<code>~a</code>	<code>~0</code>	-1	Bitwise complement of a.

Some Facts

- An n -bit bitmap may be represented with $(n + 7)/8$ bytes.
- We can represent all possible 2^n subset configurations.
- All operations take constant time.
- All operations are correct: there are no errors.

When Big Gets Huge

- Bitmaps are nice if your sets are of the form $\{0, \dots, n-1\}$.
- Many applications rely on sets U , with $n = |U|$, but
 - $n \ll \max_{u \in U}(u)$.
- We could still represent them as bitmaps.
- However, this would require $\max_{u \in U}(u)$ bits.
- Should this happen, a lot of time is wasted on swopping.
- A datastructure that fits in to memory would be *much* faster:
 - Even if the resulting “set” operations are sometimes inaccurate.

Introduction to Bloom Filters

- A *Bloom Filter* is a probabilistic version of a “set.”
- It requires much less memory than a bitmap-based set.
- It cannot answer questions about size.
- You can add to the set, but you cannot remove from the set.
- It can be used to answer set membership queries:
 - `contains(int key)?`
 - The accuracy of the answer depends on the answer:
 - `false`: Set definitely doesn't contain key.
 - `true`: Set contains key with a certain degree of confidence.

Introduction to Bloom Filters (Continued)

- Let u be a random member from the set member universe U .
- Furthermore, let B be a Bloom Filter.
- There are four scenarios.
 - B returns false and $u \notin B$: u is definitely not in B .
 - B returns false and $u \in B$: this situation cannot occur.
 - B returns true and $u \in B$: $u \in B$.
 - B returns true and $u \notin B$: *False positive*. (May happen).

Introduction to Bloom Filters (Continued)

- Let u be a random member from the set member universe U .
- Furthermore, let B be a Bloom Filter.
- There are four scenarios.
 - B returns false and $u \notin B$: u is definitely not in B .
 - B returns false and $u \in B$: this situation cannot occur.
 - B returns true and $u \in B$: $u \in B$.
 - B returns true and $u \notin B$: False positive. (May happen).

Introduction to Bloom Filters (Continued)

- Let u be a random member from the set member universe U .
- Furthermore, let B be a Bloom Filter.
- There are four scenarios.
 - B returns false and $u \notin B$: u is definitely not in B .
 - B returns false and $u \in B$: this situation cannot occur.
 - B returns true and $u \in B$: $u \in B$.
 - B returns true and $u \notin B$: *False positive*. (May happen).

Introduction to Bloom Filters (Continued)

- Let u be a random member from the set member universe U .
- Furthermore, let B be a Bloom Filter.
- There are four scenarios.
 - B returns false and $u \notin B$: u is definitely not in B .
 - B returns false and $u \in B$: this situation cannot occur.
 - B returns true and $u \in B$: $u \in B$.
 - B returns true and $u \notin B$: *False positive. (May happen).*

Introduction to Bloom Filters (Continued)

- Let u be a random member from the set member universe U .
- Furthermore, let B be a Bloom Filter.
- There are four scenarios.
 - B returns false and $u \notin B$: u is definitely not in B .
 - B returns false and $u \in B$: this situation cannot occur.
 - B returns true and $u \in B$: $u \in B$.
 - B returns true and $u \notin B$: *False positive*. (May happen).
- In general we don't know whether $u \in B$ or $u \notin B$.
 - (That's why we're using B .)

Introduction to Bloom Filters (Continued)

- Let u be a random member from the set member universe U .
- Furthermore, let B be a Bloom Filter.
- There are four scenarios.
 - B returns false and $u \notin B$: u is definitely not in B .
 - B returns false and $u \in B$: this situation cannot occur.
 - B returns true and $u \in B$: $u \in B$.
 - B returns true and $u \notin B$: *False positive*. (May happen).
- In general we don't know whether $u \in B$ or $u \notin B$.
 - (That's why we're using B .)
- We never have to worry when B returns false.
- We only have to be careful that false positives may occur.
- So when B returns true, the result may be incorrect.

Observations

- Let's revisit our bitmaps.
- We may view a bitmap as a key-value system.
- The keys and values are equal.
- With perfect hashing we could decide set membership of n members using n bits:
 - We hash each set members to its hash code.
 - The hash codes should be in the range $\{0, \dots, n-1\}$.
 - Use the hash code as an index in the array.
- Unfortunately, collisions are very common.

Implementation

- Let U be the universe of candidate set members.
- Let $n = |U|$ and let $I = \{0, \dots, m-1\}$, with $m \geq n$.
- With one *perfect* hash function, $h_0 : U \rightarrow I$ we can decide set membership with a bitmap.
- Bloom Filters decide set membership with *several* hash functions, $h_i : U \rightarrow I$.
- An empty filter is represented with m bits: each bit is 0.
- To add $u \in U$, we set Bit $h_i(u)$ to 1 for each hash function $h_i(\cdot)$.
- Then u is not contained if Bit $h_i(u)$ is 0 for some hash function.
- Otherwise, u is in the set, but false positives may occur.

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
-----	----------	----------	---------------

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110
7	2	3	1100

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110
7	2	3	1100

Keys Added Bits Set $0 \in B?$ $1 \in B?$ $7 \in B?$

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110
7	2	3	1100

Keys Added	Bits Set	$0 \in B?$	$1 \in B?$	$7 \in B?$
{ }	0000	-	-	-

Introduction

Bitmaps

Bloom Filters

Implementation

Applications

Properties

Question Time

For Monday

Summary

About this Document

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110
7	2	3	1100

Keys Added	Bits Set	$0 \in B?$	$1 \in B?$	$7 \in B?$
$\{\}$	0000	-	-	-
$\{0\}$	0011	+	-	-

Introduction

Bitmaps

Bloom Filters

Implementation

Applications

Properties

Question Time

For Monday

Summary

About this Document

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110
7	2	3	1100

Keys Added	Bits Set	$0 \in B?$	$1 \in B?$	$7 \in B?$
$\{\}$	0000	-	-	-
$\{0\}$	0011	+	-	-
$\{1\}$	0110	-	+	-

Introduction

Bitmaps

Bloom Filters

Implementation

Applications

Properties

Question Time

For Monday

Summary

About this Document

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110
7	2	3	1100

Keys Added	Bits Set	$0 \in B?$	$1 \in B?$	$7 \in B?$
$\{\}$	0000	-	-	-
$\{0\}$	0011	+	-	-
$\{1\}$	0110	-	+	-
$\{7\}$	1100	-	-	+

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110
7	2	3	1100

Keys Added	Bits Set	$0 \in B?$	$1 \in B?$	$7 \in B?$
$\{\}$	0000	-	-	-
$\{0\}$	0011	+	-	-
$\{1\}$	0110	-	+	-
$\{7\}$	1100	-	-	+
$\{0, 1\}$	0111	+	+	-

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110
7	2	3	1100

Keys Added	Bits Set	$0 \in B?$	$1 \in B?$	$7 \in B?$
$\{\}$	0000	-	-	-
$\{0\}$	0011	+	-	-
$\{1\}$	0110	-	+	-
$\{7\}$	1100	-	-	+
$\{0, 1\}$	0111	+	+	-
$\{0, 7\}$	1111	+	+	+

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110
7	2	3	1100

Keys Added	Bits Set	$0 \in B?$	$1 \in B?$	$7 \in B?$
$\{\}$	0000	-	-	-
$\{0\}$	0011	+	-	-
$\{1\}$	0110	-	+	-
$\{7\}$	1100	-	-	+
$\{0, 1\}$	0111	+	+	-
$\{0, 7\}$	1111	+	+	+
$\{1, 7\}$	1110	-	+	+

Implementation

- For example, let's assume $U = \{0, 1, 7\}$.
- Furthermore, let's use 4 bits and 2 hash functions:
 - $h_0(k) = k + 1 \bmod 3$ and $h_1(k) = k \bmod 4$.

k	$h_0(k)$	$h_1(k)$	Union of Bits
0	1	0	0011
1	2	1	0110
7	2	3	1100

Keys Added	Bits Set	$0 \in B?$	$1 \in B?$	$7 \in B?$
$\{\}$	0000	-	-	-
$\{0\}$	0011	+	-	-
$\{1\}$	0110	-	+	-
$\{7\}$	1100	-	-	+
$\{0, 1\}$	0111	+	+	-
$\{0, 7\}$	1111	+	+	+
$\{1, 7\}$	1110	-	+	+
$\{0, 1, 7\}$	1111	+	+	+

Applications: Key-Value Storage Systems

- One application of Bloom Filters are key-value storage systems.
- These systems use *slow* secondary media to store key values.
- Not all candidate keys correspond to values.
- For simplicity, let's assume there's only one slow disk.
- We want to know if some value (and if so which) has key k .
- The query $B.\text{contains}(k)$ may result in three cases:
 - 1 B returns false:
 - Avoids disk access.
 - 2 B returns true and $k \in B$.
 - Looking up the value using a slow method, returns the value.
 - 3 B returns true and $k \notin B$.
 - Looking up the value fails, we return \perp .
- If there are enough queries for bogus keys then this saves time.

Application: Detecting Malicious Websites

- Let's assume we want to know whether a given URL is malicious.
- We hash the URL into hash codes h_0, \dots, h_{n-1} .
- If Bit $h_i = 0$ for some i s.t. $0 \leq i < n$, then the URL isn't known.
- Otherwise the URL is "known" (but this may be a false positive).
- There are two possibilities:
 - We trust the answer.
 - We use a slow database operation to see if the URL is known.

Probabilistic Spell Checkers

- We start with an empty Bloom Filter, B .
- For each allowed word, we add the word's hash code to B .
- The user enters a word.
- We compute the hash codes of the word.
- If some hash code bits aren't in the filter, the word is invalid.
- Otherwise, we assume the word is valid.

Pre-Processing Database Joins

- Let's assume we have two database tables T_1 and T_2 .
- Let's compute the join $T_1 \bowtie T_2$.
- This is expensive and takes $\mathcal{O}(|T_1| \times |T_2|)$ worst-case time.
- Removing redundant rows from T_1 and T_2 takes much time.
- This also takes $\mathcal{O}(|T_1| \times |T_2|)$ worst-case time.
- Using Bloom Filters we do this in time $\mathcal{O}(|T_1| + |T_2|)$.

Pre-Processing Database Joins (Continued)

- Let's remove the redundant rows of T_1 .
- Let S be the intersection of the scopes of T_1 and T_1 .
- We start with an empty filter B .
- There are two phases:
 - 1 Add hash values of the rows in T_2 to B .
 - 2 Remove rows from T_1 whose hash values aren't in B .
- In practice, this pre-processing speeds up the join.

- Time to decide membership is independent of current “size.”
 - With k constant-time hash functions we need $\mathcal{O}(k)$ time.
- Reliability may be improved by increasing the number of bits.
- Let R be the *false positive rate*.
- For ± 4.8 extra bits per member, R reduces by a factor of ± 10 .

Questions Anybody?

For Monday

- Study the presentation,
- Implement the bitmap class from scratch.

Summary

- We've studied n -bit bitmaps and m -bit Bloom Filters.
- Bitmap decides which numbers in $\{0, \dots, n-1\}$ are in the set.
- The i -th member is in the set if and only if the i th bit is set.
- They are 100% accurate.
- With perfect hashing you can use them for other kinds of sets.
- Bloom Filters are probabilistic data structures.
- They use several hash functions, $h_i(\cdot)$.
- Candidate, c , isn't in the filter if Bit $h_i(c)$ is 0, for some i .
- Otherwise, Bit $h_i(c)$ is 1 for each i .
- In this case u is in the set, but false positives may occur.

About this Document

- This document was created with pdf \LaTeX atex.
- The \LaTeX document class is beamer.