


Network Layer - Contents

- ❑ Introduction
- ❑ Virtual circuit and datagram networks
- ❑ What's inside a router
- ❑ IP: Internet Protocol
 - Datagram format
 - IPv4 addressing
 - ICMP
 - IPv6
- ❑ Routing algorithms
 - Link state
 - Distance Vector
 - Hierarchical routing
- ❑ Routing in the Internet
 - RIP, OSPF
 - BGP
- ❑ Broadcast and multicast routing
- ❑ Software Defined Networks

University College Cork CS3506

1


Interplay between routing, forwarding



University College Cork CS3506

2

Graph abstraction



Graph: $G = (N, E)$

N = set of routers = $\{ u, v, w, x, y, z \}$

E = set of links = $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$


Remark: Graph abstraction is useful in other network contexts

Example: P2P, where N is set of peers and E is set of TCP connections

University College Cork CS3506

3

Graph abstraction: costs



• $c(x, x') = \text{cost of link } (x, x')$

- e.g., $c(w, z) = 5$

• cost could always be 1, or
inversely related to bandwidth,
or related to congestion

Cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

Question: What's the least-cost path between u and z ?

Routing algorithm: algorithm that finds least-cost path

University College Cork CS3506

4

Routing Algorithm classification

Global or decentralized information?

Global:

- ❑ all routers have complete topology, link cost info
- ❑ “link state” algorithms

Decentralized:

- ❑ router knows physically-connected neighbours, link costs to neighbours
- ❑ iterative process of computation, exchange of info with neighbours
- ❑ “distance vector” algorithms

University College Cork CS3506

Static or dynamic?

Static:

- ❑ routes change slowly over time

Dynamic:

- ❑ routes change more quickly
 - periodic update
 - in response to link cost changes

5

A Link-State Routing Algorithm

Dijkstra's algorithm

- ❑ net topology, link costs known to all nodes
 - accomplished via “link state broadcast”
 - all nodes have same info
- ❑ computes least cost paths from one node (‘source’) to all other nodes
 - gives forwarding table for that node
- ❑ iterative: after k iterations, know least cost path to k destinations

University College Cork CS3506

Notation:

- ❑ $c(x,y)$: link cost from node x to y; $= \infty$ if not direct neighbours
- ❑ $D(v)$: current value of cost of path from source to destination v
- ❑ $p(v)$: predecessor node along path from source to v
- ❑ N' : set of nodes whose least cost path definitively known

6

Dijkstra's Algorithm

```

1 Initialization:
2  N' = {u}
3  for all nodes v
4    if v adjacent to u
5      then D(v) = c(u,v)
6    else D(v) = ∞
7
8 Loop
9  find w not in N' such that D(w) is a minimum
10 add w to N'
11 update D(v) for all v adjacent to w and not in N' :
12   D(v) = min( D(v), D(w) + c(w,v) )
13   /* new cost to v is either old cost to v or known
14      shortest path cost to w plus cost from w to v */
15 until all nodes in N'


```

University College Cork CS3506

7

Dijkstra's algorithm: example

Step	N'	D(v),p(v)	D(w),p(w)	D(x),p(x)	D(y),p(y)	D(z),p(z)
0	u	2,u	5,u	1,u	∞	∞
1	ux	2,u	4,x		2,x	∞
2	uxy	2,u	3,y			4,y
3	uxyv		3,y			4,y
4	uxyvw					4,y
5	uxyvwz					




University College Cork CS3506

8

Dijkstra's algorithm: example (2)

Resulting shortest-path tree from u:



Resulting forwarding table in u:

destination	link
v	(u,v)
x	(u,x)
y	(u,x)
w	(u,x)
z	(u,x)

University College Cork CS3506

9


Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- each iteration: need to check all nodes, w , not in N'
- $n(n+1)/2$ comparisons: $O(n^2)$
- more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- e.g., link cost = amount of carried traffic



University College Cork CS3506

10

Distance Vector Algorithm

Bellman-Ford Equation (dynamic programming)

Define

$d_x(y) :=$ cost of least-cost path from x to y

Then


$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$

where min is taken over all neighbours v of x

University College Cork CS3506

11

Bellman-Ford example



Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

Node that achieves minimum is next
hop in shortest path \rightarrow forwarding table

University College Cork CS3506

12

Distance Vector Algorithm

- $D_x(y)$ = estimate of least cost from x to y
- Node x knows cost to each neighbour v :
 $c(x,v)$
- Node x maintains distance vector $D_x = [D_x(y): y \in N]$
- Node x also maintains its neighbours' distance vectors
 - For each neighbour v , x maintains $D_v = [D_v(y): y \in N]$

University College Cork CS3506

13

Distance vector algorithm (4)

Basic idea:

- From time-to-time, each node sends its own distance vector estimate to neighbours
- Asynchronous
- When a node x receives new DV estimate from neighbour, it updates its own DV using B-F equation:
$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \quad \text{for each node } y \in N$$
- Under minor, natural conditions, the estimate $D_x(y)$ converge to the actual least cost $d_x(y)$

University College Cork CS3506

14

Distance Vector Algorithm (5)

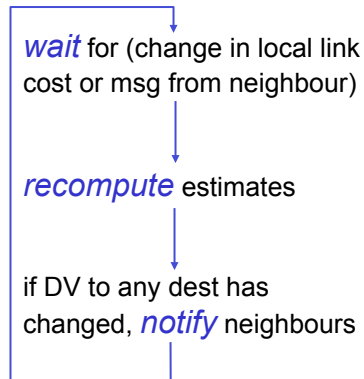
Iterative, asynchronous:
each local iteration caused by:

- local link cost change
- DV update message from neighbour

Distributed:


- each node notifies neighbours *only* when its DV changes
 - neighbours then notify their neighbours if necessary

Each node:




University College Cork CS3506

15



University College Cork CS3506

16




University College Cork CS3506

17

Distance Vector: link cost changes

Link cost changes:

- node detects local link cost change
- updates routing info, recalculates distance vector
- if DV changes, notify neighbours



“good
news
travels
fast”

At time t_0 , y detects the link-cost change, updates its DV, and informs its neighbours.

At time t_1 , z receives the update from y and updates its table. It computes a new least cost to x and sends its neighbours its DV.

At time t_2 , y receives z's update and updates its distance table. y's least costs do not change and hence y does not send any message to z.


University College Cork CS3506

18

Distance Vector: link cost changes

Link cost changes:

- ❑ good news travels fast
- ❑ bad news travels slowly - “count to infinity” problem!
- ❑ 44 iterations before algorithm stabilizes: see example in textbook



Poisoned reverse:

- ❑ If Z routes through Y to get to X :
 - Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)
- ❑ will this completely solve count to infinity problem?

University College Cork CS3506

19

Comparison of LS and DV algorithms

Message complexity

- ❑ **LS:** with n nodes, E links, $O(nE)$ msgs sent
- ❑ **DV:** exchange between neighbours only
 - convergence time varies

Speed of Convergence

- ❑ **LS:** $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- ❑ **DV:** convergence time varies
 - may be routing loops
 - count-to-infinity problem

Robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect *link* cost
- each node computes only its *own* table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - error propagate thru network

University College Cork CS3506

20

Hierarchical Routing

Our routing study thus far - idealization

- ❑ all routers identical
 - ❑ network “flat”
- ... *not* true in practice

scale: with 200 million destinations:

- ❑ can't store all dest's in routing tables!
- ❑ routing table exchange would swamp links!

administrative autonomy

- ❑ internet = network of networks
- ❑ each network admin may want to control routing in its own network


Hierarchical Routing

- ❑ aggregate routers into regions, “**autonomous systems**” (AS)
- ❑ routers in same AS run same routing protocol
 - “**intra-AS**” routing protocol
 - routers in different AS can run different intra-AS routing protocol

Gateway router

- ❑ Direct link to router in another AS

Interconnected ASes



- forwarding table configured by both intra- and inter-AS routing algorithm
 - intra-AS sets entries for internal destinations
 - inter-AS & intra-AS sets entries for external destinations

University College Cork CS3506

23


Inter-AS tasks

- suppose router in AS1 receives datagram destined outside of AS1:
 - router should forward packet to gateway router, but which one?

AS1 must:

1. learn which destinations are reachable through AS2, which through AS3
2. propagate this reachability information to all routers in AS1

Job of inter-AS routing!




University College Cork CS3506

24

Example: Setting forwarding table in router 1d

- suppose AS1 learns (via inter-AS protocol) that subnet **x** reachable via AS3 (gateway 1c) but not via AS2.
- inter-AS protocol propagates reachability info to all internal routers.
- router 1d determines from intra-AS routing info that its interface **I** is on the least cost path to 1c.
 - installs forwarding table entry (**x,I**)




University College Cork CS3506

25

Example: Choosing among multiple ASes

- now suppose AS1 learns from inter-AS protocol that subnet **x** is reachable from AS3 *and* from AS2.
- to configure forwarding table, router 1d must determine towards which gateway it should forward packets for dest **x**.
 - this is also job of inter-AS routing protocol!

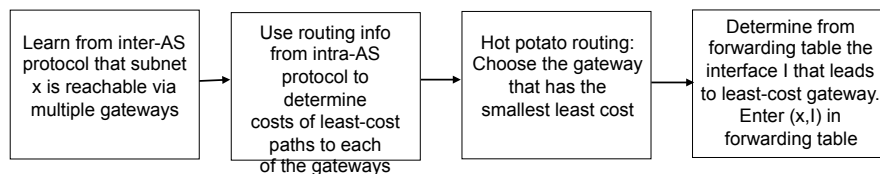


University College Cork CS3506

26

Example: Choosing among multiple ASes

- ❑ now suppose AS1 learns from inter-AS protocol that subnet **x** is reachable from AS3 *and* from AS2.
- ❑ to configure forwarding table, router 1d must determine towards which gateway it should forward packets for dest **x**.
 - this is also job of inter-AS routing protocol!
- ❑ **hot potato routing**: send packet towards closest of two routers.



University College Cork CS3506

27

Network Layer - Contents

- ❑ Introduction
- ❑ Virtual circuit and datagram networks
- ❑ What's inside a router
- ❑ IP: Internet Protocol
 - Datagram format
 - IPv4 addressing
 - ICMP
 - IPv6
- ❑ Routing algorithms
 - Link state
 - Distance Vector
 - Hierarchical routing
- ❑ Routing in the Internet
 - RIP, OSPF
 - BGP
- ❑ Broadcast and multicast routing
- ❑ Software Defined Networks

University College Cork CS3506


28

Intra-AS Routing

- ❑ also known as **Interior Gateway Protocols (IGP)**
- ❑ most common Intra-AS routing protocols:
 - RIP: Routing Information Protocol
 - OSPF: Open Shortest Path First
 - IGRP: Interior Gateway Routing Protocol (Cisco proprietary)

RIP (Routing Information Protocol)

- ❑ distance vector algorithm
- ❑ included in BSD-UNIX Distribution in 1982
- ❑ distance metric: # of hops (max = 15 hops)




From router A to subnets:

<u>destination</u>	<u>hops</u>
u	1
v	2
w	2
x	3
y	3
z	2

RIP advertisements

- distance vectors: exchanged among neighbours every 30 sec via Response Message (also called **advertisement**)
- each advertisement: list of up to 25 destination subnets within AS

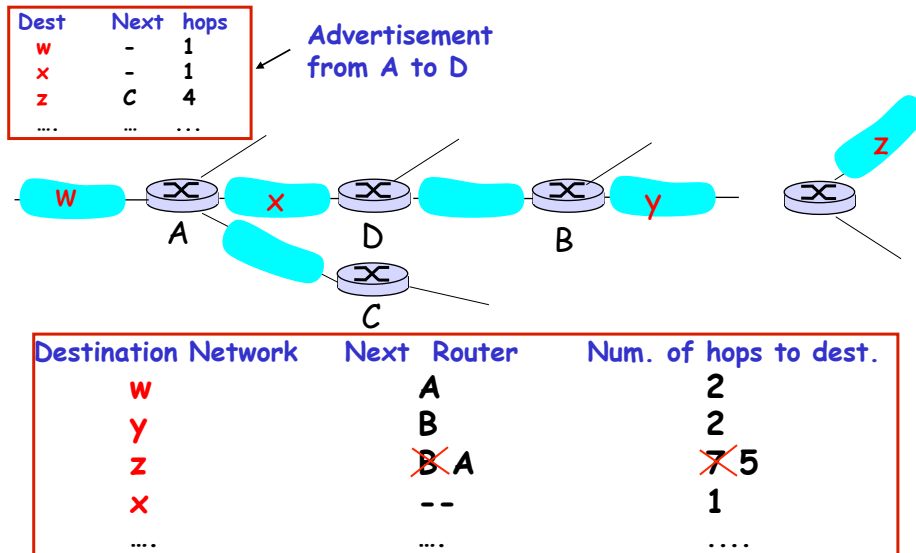
RIP: Example



Destination Network	Next Router	Num. of hops to dest.
W	A	2
Y	B	2
Z	B	7
X	--	1
...

Routing/Forwarding table in D

RIP: Example



University College Cork CS3506

Routing/Forwarding table in D

33

RIP: Link Failure and Recovery

If no advertisement heard after 180 sec --> neighbour/
link declared dead


- routes via neighbour invalidated
- new advertisements sent to neighbours
- neighbours in turn send out new advertisements (if tables changed)
- link failure info quickly (?) propagates to entire net
- *poison reverse* used to prevent ping-pong loops (infinite distance = 16 hops)

University College Cork CS3506

34

RIP Table processing

- ❑ RIP routing tables managed by **application-level** process called route-d (daemon)
- ❑ advertisements sent in UDP packets, periodically repeated



University College Cork CS3506

35

OSPF (Open Shortest Path First)

- ❑ “open”: publicly available
- ❑ uses Link State algorithm
 - LS packet dissemination
 - topology map at each node
 - route computation using Dijkstra’s algorithm
- ❑ OSPF advertisement carries one entry per neighbour router
- ❑ advertisements disseminated to **entire** AS (via flooding)
 - carried in OSPF messages directly over IP (rather than TCP or UDP)

University College Cork CS3506

36


OSPF “advanced” features (not in RIP)

- ❑ **security**: all OSPF messages authenticated (to prevent malicious intrusion)
- ❑ **multiple** same-cost **paths** allowed (only one path in RIP)
- ❑ For each link, multiple cost metrics for different **TOS** (e.g., satellite link cost set “low” for best effort; high for real time)
- ❑ integrated uni- and **multicast** support:
 - Multicast OSPF (MOSPF) uses same topology data base as OSPF
- ❑ **hierarchical** OSPF in large domains.

University College Cork CS3506

37

Hierarchical OSPF



University College Cork CS3506

38

Hierarchical OSPF

- ❑ **two-level hierarchy:** local area, backbone.
 - Link-state advertisements only in area
 - each nodes has detailed area topology; only know direction (shortest path) to nets in other areas.
- ❑ **area border routers:** “summarize” distances to nets in own area, advertise to other Area Border routers.
- ❑ **backbone routers:** run OSPF routing limited to backbone.
- ❑ **boundary routers:** connect to other AS' s.

University College Cork CS3506

39

Internet inter-AS routing: BGP


- ❑ **BGP (Border Gateway Protocol):** *the de facto standard*
- ❑ BGP provides each AS a means to:
 1. Obtain subnet reachability information from neighbouring ASs.
 2. Propagate reachability information to all AS-internal routers.
 3. Determine “good” routes to subnets based on reachability information and policy.
- ❑ allows subnet to advertise its existence to rest of Internet: *“I am here”*

University College Cork CS3506

40

BGP basics

- pairs of routers (BGP peers) exchange routing info over semi-permanent TCP connections: **BGP sessions**
 - BGP sessions need not correspond to physical links.
- when AS2 advertises a prefix to AS1:
 - AS2 **promises** it will forward datagrams towards that prefix.
 - AS2 can aggregate prefixes in its advertisement




University College Cork CS3506

41

Distributing reachability info

- using eBGP session between 3a and 1c, AS3 sends prefix reachability info to AS1.
 - 1c can then use iBGP to distribute new prefix info to all routers in AS1
 - 1b can then re-advertise new reachability info to AS2 over 1b-to-2a eBGP session
- when router learns of new prefix, it creates entry for prefix in its forwarding table.



University College Cork CS3506

42

Path attributes & BGP routes

- ❑ advertised prefix includes BGP attributes.
 - prefix + attributes = “route”
- ❑ two important attributes:
 - **AS-PATH**: contains ASs through which prefix advertisement has passed: e.g, AS 67, AS 17
 - **NEXT-HOP**: indicates specific internal-AS router to next-hop AS. (may be multiple links from current AS to next-hop-AS)
- ❑ when gateway router receives route advertisement, uses **import policy** to accept/decline.

University College Cork CS3506

43

BGP route selection

- ❑ router may learn about more than 1 route to some prefix. Router must select route.
- ❑ elimination rules:
 1. local preference value attribute: policy decision
 2. shortest AS-PATH
 3. closest NEXT-HOP router: hot potato routing
 4. additional criteria

University College Cork CS3506

44


BGP messages

- ❑ BGP messages exchanged using TCP.
- ❑ BGP messages:
 - **OPEN**: opens TCP connection to peer and authenticates sender
 - **UPDATE**: advertises new path (or withdraws old)
 - **KEEPALIVE** keeps connection alive in absence of UPDATES; also ACKs OPEN request
 - **NOTIFICATION**: reports errors in previous msg; also used to close connection

University College Cork CS3506

45

BGP routing policy




- ❑ A,B,C are **provider networks**
- ❑ X,W,Y are customer (of provider networks)
- ❑ X is **dual-homed**: attached to two networks
 - X does not want to route from B via X to C
 - .. so X will not advertise to B a route to C

University College Cork CS3506

46

BGP routing policy (2)



- ❑ A advertises path *AW* to B
- ❑ B advertises path *BAW* to X
- ❑ Should B advertise path *BAW* to C?
 - No way! B gets no “revenue” for routing *CBAW* since neither *W* nor *C* are B’s customers
 - B wants to force *C* to route to *w* via *A*
 - B wants to route *only* to/from its customers!

University College Cork CS3506

47

Why different Intra- and Inter-AS routing ?

Policy:

- ❑ Inter-AS: admin wants control over how its traffic routed, who routes through its net.
- ❑ Intra-AS: single admin, so no policy decisions needed

Scale:

- ❑ hierarchical routing saves table size, reduced update traffic

Performance:

- ❑ Intra-AS: can focus on performance
- ❑ Inter-AS: policy may dominate over performance

University College Cork CS3506

48

Network Layer - Contents


- ❑ Introduction
- ❑ Virtual circuit and datagram networks
- ❑ What's inside a router
- ❑ IP: Internet Protocol
 - Datagram format
 - IPv4 addressing
 - ICMP
 - IPv6
- ❑ Routing algorithms
 - Link state
 - Distance Vector
 - Hierarchical routing
- ❑ Routing in the Internet
 - RIP, OSPF
 - BGP
- ❑ Broadcast and multicast routing
- ❑ Software Defined Networks

University College Cork CS3506

49

Broadcast Routing

- ❑ deliver packets from source to all other nodes (contrast with unicast routing)
- ❑ Could use source duplication, but it's inefficient:



- ❑ Another problem for source duplication: how does source determine recipient addresses?

University College Cork CS3506

50

In-network duplication


- ❑ flooding: when node receives brdcst pkt, sends copy to all neighbours
 - Problems: cycles & broadcast storm
- ❑ controlled flooding: node only brdcsts pkt if it hasn't brdcst same packet before
 - Node keeps track of pkt ids already brdcsted
 - Or reverse path forwarding (RPF): only forward pkt if it arrived on shortest path between node and source
- ❑ Or use a spanning tree
 - No redundant packets received by any node

University College Cork CS3506


51

Spanning Tree

- ❑ First construct a spanning tree
- ❑ Nodes forward copies only along spanning tree



(a) Broadcast initiated at A




(b) Broadcast initiated at D

University College Cork CS3506


52

Spanning Tree: Creation

- ❑ Center node
- ❑ Each node sends unicast join message to center node
 - Message forwarded until it arrives at a node already belonging to spanning tree



(a) Stepwise construction of spanning tree



(b) Constructed spanning tree

University College Cork CS3506

53

Multicast Routing


- ❑ Unlike broadcast routing, the goal of multicast is to deliver a packet to a defined *subset* of the network hosts
- ❑ Ideal for live events such as concerts, live TV, etc
 - A subset of network hosts are interested
 - Sending a unicast to each interested host doesn't scale
 - Each packet must be sent once for each interested host
 - Same content traverses a link multiple times

University College Cork CS3506

54

Multicast Routing: Problem Statement

- **Goal:** find a tree (or trees) connecting routers having local mcast group members
 - **tree:** not all paths between routers used
 - **source-based:** different tree from each sender to rcvrs
 - **shared-tree:** same tree used by all group members



University College Cork CS3506

55

Approaches for building mcast trees

- **source-based tree:** one tree per source
 - shortest path trees
 - reverse path forwarding
- **group-shared tree:** group uses one tree
 - minimal spanning tree (known as a Steiner tree)
 - tree based at center node (i.e. router)
- A variety of protocols in use including
 - DVMRP - distance vector multicast
 - PIM - Protocol Independent Multicast

University College Cork CS3506

56

Network Layer - Contents

- ❑ Introduction
- ❑ Virtual circuit and datagram networks
- ❑ What's inside a router
- ❑ IP: Internet Protocol
 - Datagram format
 - IPv4 addressing
 - ICMP
 - IPv6
- ❑ Routing algorithms
 - Link state
 - Distance Vector
 - Hierarchical routing
- ❑ Routing in the Internet
 - RIP, OSPF
 - BGP
- ❑ Broadcast and multicast routing
- ❑ Software Defined Networks

University College Cork CS3506

57

Software-Defined Networks


- ❑ Distributed implementation of network control (eg. Routing, NAT, firewalls)
- ❑ Difficult to make changes to network behaviour
 - Features implemented in software provided by router/switch vendors
 - Long delays waiting for protocol standardisation
 - Stifles innovation
- ❑ Difficult to manage large networks with many routers/switches

University College Cork CS3506

58

Today's Internet: Data Plane

Data plane:
Packet
streaming



Forward, filter, buffer, mark,
rate-limit, and measure packets


SDN slides based on originals by Jennifer Rexford @ Princeton

University College Cork CS3506

59

Today's Internet: Control Plane

Control plane:
Distributed algorithms




Track topology changes, compute
routes, install forwarding rules

University College Cork CS3506

60

Today's Internet: Management

Management plane:
Human time scale




Collect measurements and configure
the equipment

University College Cork CS3506

61

Software-Defined Networking



University College Cork CS3506

62

Data-Plane: Simple Packet Handling

❑ Simple packet-handling rules

- Pattern: match packet header bits
- Actions: drop, forward, modify, send to controller
- Priority: disambiguate overlapping patterns
- Counters: #bytes and #packets



1. src=1.2.*.*, dest=3.4.5.* → drop
2. src = *.*.*.*, dest=3.4.*.* → forward(2)
3. src=10.1.2.3, dest=*.*.*.* → send to controller

University College Cork CS3506

63

Unifies Different Core Devices

❑ Router

- Match: longest destination IP prefix
- Action: forward out a link

❑ Switch

- Match: destination MAC address
- Action: forward or flood

❑ Firewall

- Match: IP addresses and TCP/UDP port numbers
- Action: permit or deny

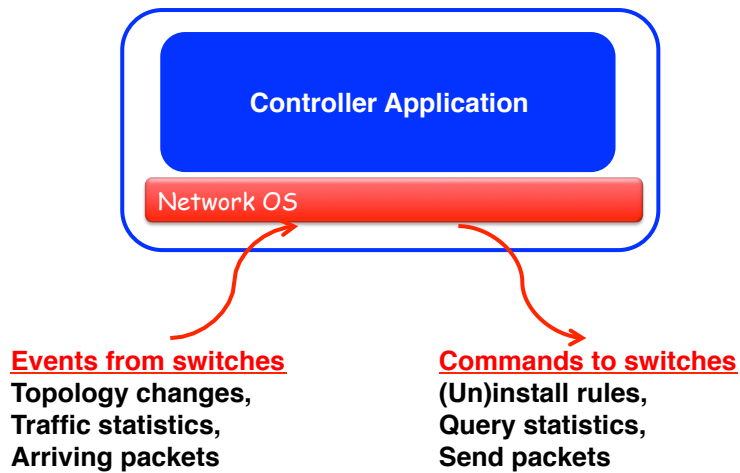
❑ NAT

- Match: IP address and port
- Action: rewrite address and port

University College Cork CS3506

64

Programmable Net. Controllers



University College Cork CS3506

65

Example Applications


- ☐ **Dynamic access control**
- ☐ Seamless mobility/migration
- ☐ **Server load balancing**
- ☐ Network virtualization
- ☐ Using multiple wireless access points
- ☐ Energy-efficient networking
- ☐ Adaptive traffic monitoring
- ☐ Denial-of-Service attack detection

University College Cork CS3506

66

Dynamic Access Control

- ☐ Inspect first packet of a connection
- ☐ Consult the access control policy
- ☐ Install rules to block or route traffic

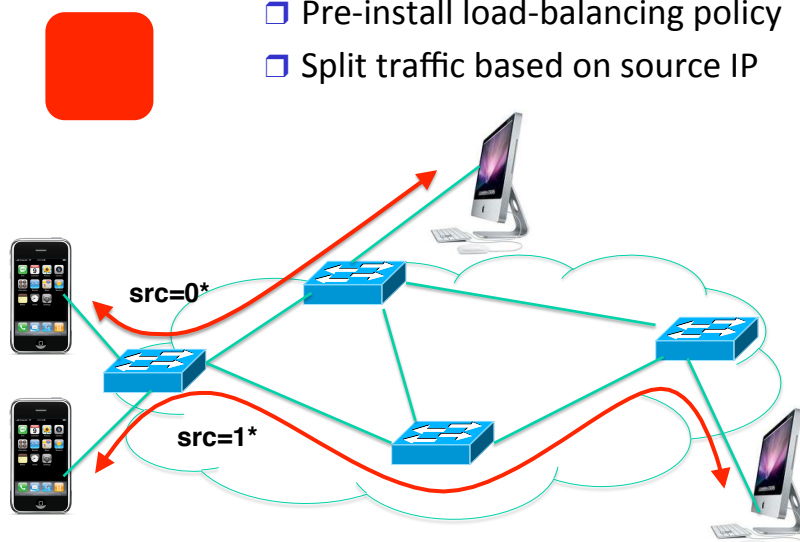


University College Cork CS3506

67

Server Load Balancing

- ☐ Pre-install load-balancing policy
- ☐ Split traffic based on source IP




University College Cork CS3506

68

Controller Delays

- ❑ Controller is much slower than the switch
- ❑ Processing packets leads to delay and overhead
- ❑ Need to keep most packets in the “fast path”



University College Cork CS3506

69

Programming

- ❑ SDN makes programming possible
 - Network-wide view at controller
 - Direct control over data plane
 - Need for high-level abstractions
 - ❑ Plenty of room for bugs
 - Still a complex, distributed system, concurrency issues
 - ❑ Need for testing techniques
 - Controller applications
 - Controller and switches
 - Rules installed in the switches
- SDN has significant industry momentum*

University College Cork CS3506

70

Summary

- ❑ Introduction
- ❑ Virtual circuit and datagram networks
- ❑ What's inside a router
- ❑ IP: Internet Protocol
 - Datagram format
 - IPv4 addressing
 - ICMP
 - IPv6
- ❑ Routing algorithms
 - Link state
 - Distance Vector
 - Hierarchical routing
- ❑ Routing in the Internet
 - RIP, OSPF
 - BGP
- ❑ Broadcast and multicast routing
- ❑ Software Defined Networks