

CS4618: Artificial Intelligence I

Reinforcement Learning

Derek Bridge

School of Computer Science and Information Technology
University College Cork

Initialization


```
In [1]: %reload_ext autoreload  
        %autoreload 2  
        %matplotlib inline
```

```
In [2]: import pandas as pd  
        import numpy as np  
        import matplotlib.pyplot as plt
```

Reinforcement learning

- The agent carries out an action
- A teacher or the environment provides a **reward** (or punishment), often delayed, that acts as positive (or negative) reinforcement
 - making it more (or less) likely that the agent will execute that action if it find itself in the same or similar situation in the future
- For simplicity, in this lecture, we assume a fully-observable, deterministic environment

Reward



$$s_0 \xrightarrow[r_0]{a_0} s_1 \xrightarrow[r_1]{a_1} s_2 \xrightarrow[r_2]{a_2} \dots$$

Cumulative reward

- Cumulative reward:

$$r_0 + \gamma r_1 + \gamma^2 r_2 + \dots$$

or

$$\sum_{t=0}^{t=\infty} \gamma^t r_t$$

where γ is the **discount rate** ($0 \leq \gamma \leq 1$)

- The task of the agent is to learn an action function that maximises cumulative reward

Action-value function

- Assume 2 Boolean sensors and 3 actions
- Compare

Percept	Action
00	MOVE
01	TURN(RIGHT, 2)
10	MOVE
11	TURN(LEFT, 2)

Percept	Action	Q
00	MOVE	...
00	TURN(RIGHT, 2)	...
00	TURN(LEFT, 2)	...
01	MOVE	...
01	TURN(RIGHT, 2)	...
01	TURN(LEFT, 2)	...
10	MOVE	...
10	TURN(RIGHT, 2)	...
10	TURN(LEFT, 2)	...
11	MOVE	...
11	TURN(RIGHT, 2)	...
11	TURN(LEFT, 2)	...

- Class exercise: Suppose the agent has m touch sensors (returning 0 or 1) and n different actions. How many rows will the table contain?

What is Q ?

- $Q(\mathbf{s}, a)$ is an *estimate* of the cumulative reward the agent will receive if, having sensed \mathbf{s} , it chooses to execute action a
- Hence, having sensed \mathbf{s} , choose action a for which $Q(\mathbf{s}, a)$ is highest:

$$\arg \max_a Q(\mathbf{s}, a)$$

Class exercise

- Given this table

Percept	Action	Q
⋮	⋮	⋮
01	MOVE	0.2
01	TURN(RIGHT, 2)	0.1
01	TURN(LEFT, 2)	0.7
⋮	⋮	⋮

- Suppose \mathbf{s} is 01
- What is $\arg \max_a Q(\mathbf{s}, a)$?

Q-learning

- Start with random Q -values (or all zero)
- Improve by trial-and-error: choose actions, get rewards, update Q -values

QLearning(ϵ)

- $s = SENSE()$;
- do forever
 - $rand$ = a randomly-generated number in $[0, 1)$;
 - if $rand < \epsilon$
 - Choose action a randomly;
 - else
 - $a = \arg \max_a Q(s, a)$;
 - $r = EXECUTE(a)$;
 - $s' = SENSE()$;
 - $Q(s, a) = r + \gamma \times \max_{a'} Q(s', a')$;
 - $s = s'$;

Exploration vs. Exploitation

- Exploration:
 - Choose an action which may not be the best action according to the current Q values. But it may gain you new experience and improve the Q values
- Exploitation:
 - Choose the action which is best according to the current Q values. It may gain you reward
- The so-called ϵ -**greedy policy** (where $0 \leq \epsilon \leq 1$) is the simplest way to balance exploration and exploitation

Updating Q -values

- From the algorithm:

$$Q(s, a) = r + \gamma \times \max_{a'} Q(s', a')$$

- The new value is the reward for the latest action r plus our highest current estimate of the cumulative reward it can receive
- Over the course of repeated actions, the Q -values will get better and better:
 - When one Q -value improves then the Q -values of its immediate predecessors will also improve next time they get updated

Class exercise

- Given this table

Percept	Action	Q
:	:	:
10	MOVE	5
10	TURN(RIGHT, 2)	4
10	TURN(LEFT, 2)	1
11	MOVE	0
11	TURN(RIGHT, 2)	4
11	TURN(LEFT, 2)	6

- Suppose current percept is 10
Assuming exploitation, which action will be chosen?
- Suppose reward is 3, next percept is 11 and γ is 1
Using $Q(s, a) = r + \gamma \times \max_{a'} Q(s', a')$, update the table

Concluding remarks

- Reinforcement Learning underpins a lot of current success in game playing
- But it is seeing real use in other areas, e.g. robot motion control
- For real use, you need more sophisticated algorithms:
 - To handle non-deterministic environments
 - To improve convergence
 - To build a model of the environment
 - To scale up, and
 - To represent the policy in a way that allows the agent to generalise from what it learns

In []: